


Searching Algorithms

Er. Rudra Nepal

Searching


- Step in Problem Solving
- Searching is Performed through the State Space
- Searching accomplished by constructing a search tree
- The root of the search tree is search node



Searching

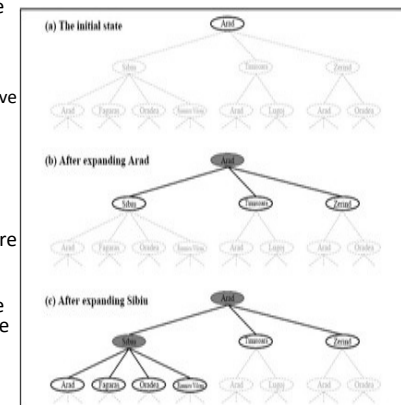
- There are many ways to represent nodes, but we will assume that a node is a data structure with five components:

- State: the state in the state space to which the node corresponds
- Parent-Node: the node in the search tree that generated this node.
- Action: the action that was applied to the parent to generate the node.
- Path-cost: the cost of the path from initial state to the node
- Depth: the number of steps along the path from initial state.



Searching: Steps

- Check whether the current state is the goal state or not
- Expand the current state to generate the new sets of states
 - The collection of nodes that have been generated but not yet expanded is called fringe.
 - Each element of fringe is a leaf node-a node with no successor in the tree.
- Choose one of the new states generated for search which entirely depend on the selected search strategy
- Repeat the above steps until the goal state is reached or there are no more states to be expanded



Searching: Criteria to Measure Performance

- The output of a search algorithm is either failure or a solution.
- Some algorithm get stuck in an infinite loop and never return an output.
- We can evaluate the algorithm's performance in four ways:
 - Completeness: Is the algorithm guaranteed to find a solution when there is one?
 - Optimality: Does the strategy find the optimal solution?
 - Time Complexity: How long does it take to find a solution?
 - Space Complexity: How much memory is needed to perform the search?
- We can express the algorithm's complexity in terms of three quantities:
 - b , the branching factor or maximum number of successor of any node
 - d , the depth of the shallowest goal node
 - m , the maximum length of any path in the state space

Searching: Types


- Uninformed Search or Blind Search
- Informed Search or Heuristic Search

Uninformed Search

- Search provided with problem definition only and no additional information about the state space
- Expansion of current state to new set of states is possible
- It can only distinguish between goal state and non-goal state
- Less effective compared to Informed search
- The uninformed search strategies are distinguished by the order in which nodes are expanded
- Various uninformed search techniques/strategies are:
 - Breadth-first Search
 - Uniform-cost Search
 - Depth-first Search
 - Depth-limited Search
 - Iterative deepening depth-first search

Breadth-first Search

- Root node is expanded first
- Then all the successors of the root node are expanded
- Then their successors are expanded and so on.
- Nodes, which are visited first will be expanded first (FIFO)
- All the nodes of depth ' d ' are expanded before expanding any node of depth ' $d+1$ '



Breadth First Search: Four Criteria

- Completeness
 - This search strategy finds the shallowest goal first
 - Complete, if the shallowest goal is at some finite depth
- Optimality
 - The shallowest goal node is not necessarily the optimal one
 - Optimal, if the path cost is a non-decreasing function of the path of the node (For example: when all the actions have the same cost)

Breadth First Search: Four Criteria

Time Complexity

- For a search tree a branching factor 'b' expanding the root yields 'b' nodes at the first level.
- Expanding 'b' nodes at first level yields b^2 nodes at the second level.
- Similarly, expanding the nodes at d th level yields b^{d+1} node at $(d+1)$ th level
- If the goal is in d th level, in the worst case, the goal node would be the last node in the d th level

Hence, We should expand 1) nodes in the $(d+1)$ level (Except the goal node itself which doesn't need to be expanded)


- So, number of nodes generated at $(d+1)$ th level = $b(b^d - 1) = b^{d+1} - b$
- Again, Total number of nodes generated = $1 + b + b^2 + \dots + b^d - b = O(b^{d+1})$
- Hence, time complexity is $O(b^{d+1})$ where, b = branching factor and d = level of goal node in the search table

Breadth First Search: Four Criteria

- Space Complexity
 - Same as time complexity
 - i.e. $O(b^{d+1})$
 - Since each node has to be kept in the memory
- Disadvantages
 - Memory Wastage
 - Irrelevant Operations
 - Time Intensive It
 - doesn't assure the optimal cost solution


Uniform Cost Search

- It expands the lowest cost node on the fringe
- The first solution is guaranteed to be the cheapest one because a cheaper one would have expanded earlier and so would have been found first
- If all step costs are equal, this is identical to breadth first search
- Required Condition: A to H
 - $ABEH=21$, $ACH=5$, $ACFH=7$, $ADIH=6$



Uniform Cost Search

- Solution: Required Operation
 - Expand A → Yield B, C, D
 - With $AB=1$, $AC=4$, $AD=3$
 - Expand B → Yield E with $ABE=13$
 - As $ABE > AC$ and $ABE > AD$
 - Expand D → Yield I with $ADI=5$
 - As $ADI > AC$
 - Expand C → Yield H and F with $ACH=5$ and $ACF=6$
 - Solution Achieved
- If all step costs are equal, it is identical breadth first search



Uniform Cost Search


- Disadvantages
 - Doesn't care about the number of steps a path has but only about their cost
 - It might get stuck in an infinite loop if it expands a node that has a zero cost action leading back to same state

Uniform Cost Search: Four Criteria

- Completeness
 - Complete, if the cost of every step is greater than or equal to some small positive constant ϵ
- Optimality
 - The same ensures optimality
- Time Complexity
 - $O(b C^*/\epsilon)$
 - Where C^* = cost of optimal path and ϵ = small positive constant
 - This complexity is much greater than that of Breadth first search
- Space Complexity
 - $O(b C^*/\epsilon)$

Depth-first Search

- Always expands the deepest node in the current fringe of the search tree
- The search proceeds immediately to the deepest level of the search tree, where the nodes have no successors (dead end)
- When a dead end is reached, the search backup to the next shallowest node that still has unexplored successors
- This strategy can be implemented by tree search with a last-in-first-out (LIFO) queue, also known as stack.



Depth First Search: Four Criteria

- Completeness
 - Can get stuck going down the wrong path when a different choice would lead to a solution near the root of the search tree
 - Not complete
- Optimality
 - The strategy might return a solution path that is longer than the optimal solution, if it starts with an unlucky path
 - Not optimal

Depth First Search: Four Criteria

- Space Complexity
 - It needs to store a single path from root to a leaf node and the remaining unexpanded sibling nodes for each node in the path
 - Once a node has been expanded, it can be removed from memory as soon as all its decedents have been fully explored
 - For a search tree of branching factor 'b' and maximum tree depth 'm', only the storage of $bm+1$ node is required
 - Hence,
 - Space Complexity = $O(bm+1)$
= $O(bm)$
- Time Complexity
 - in the worst case all the bnodes of the search tree would be generated
 - Hence,
Time Complexity= $O(bm)$

Depth-limited Search

- Modification of depth first search
- Depth first search with predetermined limit 'l'
- After the nodes at the level 'l' are explored, the search backtracks without going further deep
- Hence, it solves the infinite path problem of the depth first search strategy
- Completeness: Complete except at additional source of incompleteness if $l < d$, i.e. shallowest goal is beyond the depth limit
- Optimality: Optimal except at $l > d$
- Time Complexity = $O(b^l)$
- Space Complexity = $O(b^l)$

Iterative deepening depth first Search

- Finds the best limit by gradually increasing depth limit first to 0, then to 1, 2 and so on-until the goal is found
- Combines the benefits of the depth first and breadth first search
- In depth-limited search, the complex part is to choose good depth limit
- This strategy addresses the issue of good depth limit by trying all possible depth limits
- The process is repeated until goal is found at depth limit 'd' which is the depth of shallowest goal
- Completeness: as of Breadth First Search i.e. Complete if branching factor is finite
- Optimality: as of Breadth First Search i.e. optimal if the path cost is non decreasing function of depth
- Space Complexity = $O(bd)$
- Time Complexity = $O(bd)$

Iterative deepening depth first Search




Figure: Four iterations of iterative deepening search on a binary tree.

Comparing uninformed search strategies

Criterion	Breadth-first Search	Uniform-cost Search	Depth-first Search	Depth-limited Search	Iterative deepening Search
Completeness	Yes, if b is finite	Yes, if step cost $\geq \epsilon$ for positive ϵ	No	No	Yes, if b is finite
Optimality	Yes, if step cost are identical	Yes	No	No	Yes, if step cost are identical
Time Complexity	$O(bd+1)$	$O(bC^*/\epsilon)$	$O(bm)$	$O(bl)$	$O(bd)$
Space Complexity	$O(bd+1)$	$O(bC^*/\epsilon)$	$O(bm)$	$O(bl)$	$O(bd)$

b is branching factor


d is the depth of the shallowest solution

m is the maximum depth of search tree

l is the depth limit

Informed (Heuristic) Search Strategies

- Strategy of problem solving where problem specific knowledge is known along with problem definition
- These search find solutions more efficiently by the use of heuristics
- Heuristic is a search technique that improves the efficiency of the search process
- By eliminating the unpromising states and their descendants from consideration, heuristic algorithms can find acceptable solutions



Informed (Heuristic) Search Strategies

- Heuristics are fallible i.e. they are likely to make mistakes as well
- It is the approach following an informed guess of next step to be taken
- It is often based on experience or intuition
- Heuristic have limited information and hence can lead to suboptimal solution or even fail to find any solution at all

Best First Search

- A node is selected for expansion based on evaluation function $f(n)$
- A node with lowest evaluation function is expanded first
- The measure i.e. evaluation function must incorporate some estimate of the cost of the path from a state to the closest goal state
- The algorithm may have different evaluation function, one of such important function is the heuristic function $h(n)$
 - $h(n)$ = the estimated cost of the cheapest path from node n to the goal
 - $h(n) = 0$, if n is goal node
- Types of best first search
 - Greedy Best First Search
 - A* Search (pronounced "A-star search")

Greedy Best First Search

- The node whose state is judged to be the closest to the goal state is expanded first
- At each step it tries to be as close to the goal as it can
- It evaluates the nodes by using heuristic function hence, $f(n) = h(n)$ where, $h(n) = 0$, for the goal node
- This search resembles depth first search in the way that it prefers to follow a single path all the way to the goal or if not found till the dead end and returns back up


Greedy Best First Search: Example

Let us see how this works for route-finding problem,

- using the straight-line distance heuristic

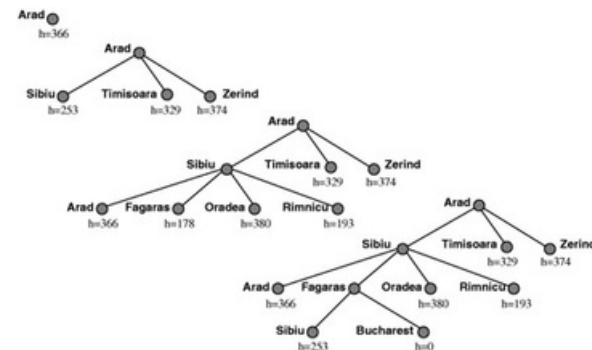
If goal is Bucharest, we will need to know the straight-

- line distance to Bucharest as shown in figure



Greedy Best First Search: Example

- Figure below shows the progress of greedy best first search using straight-line distance to find the path from Arad to Bucharest



Greedy Best First Search: Four Criteria

- Completeness
 - Can start down an infinite path and never return to any possibilities
 - Not complete
- Optimality
 - Looks for immediate best choice and doesn't make careful analysis of long term options
 - May give longer solution even if shorter solution exists
 - Not optimal
- Space Complexity
 - $O(bm)$ where, m is the maximum depth of search space, since all nodes have to be kept in memory
- Time Complexity
 - $O(bm)$

A* Search


- It evaluates nodes by combining $g(n)$, the cost to reach the node, and $h(n)$, the cost to get from node to the goal

$$f(n) = g(n) + h(n)$$


Since $g(n)$ gives the path cost from the start node to node n , and $h(n)$ is the estimated cost of the cheapest path from n to goal node, we have

$$f(n) = \text{estimated cost of the cheapest solution through } n$$
- Admissible Heuristic: $h(n)$ is admissible if it never overestimates the cost to reach the solution
 - example: $\sqrt{2}LD$ (straight line distance)
 - SLD is admissible because the shortest path between any two point is straight line

A* Search: Example



A* Search: Example




A* Search: Four Criteria

- Optimality
 - Optimal if $h(n)$ is admissible
- Space Complexity
 - $O(b)^d$ if $h(n)$ is admissible
- Completeness
 - Complete if $h(n)$ is admissible
- Time Complexity
 - $O(b)^d$ if $h(n)$ is admissible

Adversarial Search Techniques

- A game can be defined as a kind of search problem (game tree) with the following components:
 - Initial State identifying the initial position in the game and identification of the first player
 - Actions returns the set of legal moves in a state
 - Successor Function returning a list of (move, state) pairs
 - Terminal Test which is true if game is over and false otherwise. States where the game has ended are called terminal states
 - Utility function which gives a numeric value for the terminal states. Example: in TTT Lose, draw and win with -1, 0 and +1

Game Tree Example: Tic-Tac-Toe



Adversarial Search Techniques

- Minimax Algorithm
- Alpha-Beta Pruning

Minimax Algorithm

Max is considered as the first player in the game and Min as the second player

- This algorithm computes the minimax decision from the current state
- At each MAX node, pick the move with maximum utility.
- At each MIN node, pick the move with minimum utility
- It uses a recursive computation of minimax values (minimax value of a node is the utility of being in the corresponding state) of each successor state
- directly implementing some defined function


The recursion proceeds from the initial node to all the leaf nodes

- Then the minimax values are backed up through the tree as the recursion unwinds
- It performs the depth first exploration of a game tree in a complete way
- If the maximum depth of
- \square the tree is m and there are b legal moves at each point then for minimax algorithm:
 - Time Complexity=
 - Space Complexity=

$O(b^m)$

Minimax Algorithm: Computation

- In the figure, the algorithm first recurses down to the three bottom-leaf nodes and uses Utility function to discover that their values are 3, 12 and 8 respectively.
- Then it takes the minimum of these values, 3, and return it as backed-up value of node B.
- A similar process gives backed-up value of 2 for C and 2 for D.
- Finally, we take maximum of 3, 2, 2 to get backed-up value of 3 for the root node A.



Alpha-Beta Pruning

- The main disadvantage of the minimax algorithm is that all the nodes in the game tree cutoff to a certain depth are examined.
- Alpha-beta pruning helps reduce the number of nodes explored.
- When applied to a standard minimax tree, alpha beta pruning returns the same move as minimax would, but prunes away the branches which couldn't possibly influence the final decision.

Alpha-Beta Pruning

- In the figure alongside, if m is better than n for Player, we will never get n in play
- Alpha-Beta pruning gets its name from the following two parameters that describe bounds on the backed-up values:
 - α = the value of the best (i.e. highest value) choice we have found so far at any choice point along the path for MAX
 - β = the value of the best (i.e. lowest value) choice that we have found so far at any choice point along the path for MIN





Figure: General case for alpha-beta pruning

Alpha-Beta Pruning: Example



Alpha-Beta Pruning: Example

- a: The first leaf below B has the value 3. Hence, B, which is a MIN node has a value of at most 3.
- b: The second leaf below B has a value 12; MIN would avoid this move, so the value of B is still at most 3.
- c: The third leaf below B has a value of 8; we have seen all B's successor states, so the value of B is exactly 3. Now we can infer that the value of the root is at least 3, because MAX has a choice worth 3 at the node.
- d: The first leaf below C has the value 2. Hence, C, which is a MIN node, has a value of at most 2. But we know that B is worth 3, so MAX would never choose C. Therefore, there is no point in looking at the other successor states of C. This is an example of Alpha-Beta Pruning.
- e: The first leaf below D has the value 14, so D is worth at most 14. This is still higher than MAX's best alternative (i.e. 3), so we need to keep exploring D's successor states.
- f: The second successor of D is worth 5, so again we need to keep exploring. The third successor is worth 2, so now D is worth exactly 2. MAX's decision at the root is to move to B, giving a value of 3.